

The ORC and parquet files are much smaller than the CSV file. Over 10x smaller.

```
bitnami@debian:~$ hdfs dfs -du -h /user/hive/warehouse
54.2 M  54.2 M  /user/hive/warehouse/q1_table
1.6 M   1.6 M  /user/hive/warehouse/q1_table_orc
3.1 M   3.1 M  /user/hive/warehouse/q1_table_parquet
```

```
[bitnami@debian:/tmp$ sed '1d' feb.csv > febnh.csv
[bitnami@debian:/tmp$ sed '1d' mar.csv > marnh.csv
[bitnami@debian:/tmp$ head jan.csv
"YEAR","MONTH","DAY_OF_MONTH","ORIGIN","DEST","DEP_DEL15","ARR_DEL15",
2019,1,16,"BOS","PHL",0.00,0.00,
2019,1,17,"BOS","PHL",0.00,0.00,
2019,1,18,"BOS","PHL",0.00,0.00,
2019,1,19,"BOS","PHL",0.00,0.00,
2019,1,20,"BOS","PHL",,,
2019,1,21,"BOS","PHL",,,
2019,1,22,"BOS","PHL",,,
2019,1,23,"BOS","PHL",0.00,0.00,
2019,1,24,"BOS","PHL",0.00,0.00,
[bitnami@debian:/tmp$ head febnh.csv
2019,2,24,"BNA","DCA",0.00,0.00,
2019,2,24,"BNA","DCA",1.00,1.00,
2019,2,24,"BNA","DCA",1.00,1.00,
2019,2,24,"BNA","DCA",0.00,0.00,
2019,2,24,"BNA","DEN",0.00,0.00,
2019,2,24,"BNA","DEN",0.00,0.00,
2019,2,24,"BNA","DEN",0.00,0.00,
2019,2,24,"BNA","DEN",0.00,0.00,
2019,2,24,"BNA","DTW",0.00,0.00,
2019,2,24,"BNA","DTW",,,
[bitnami@debian:/tmp$ head marnh.csv
2019,3,2,"ATL","MCI",0.00,0.00,
2019,3,2,"ATL","MCI",0.00,0.00,
2019,3,2,"ATL","MCO",1.00,1.00,
2019,3,2,"ATL","MCO",0.00,0.00,
2019,3,2,"ATL","MCO",0.00,0.00,
2019,3,2,"ATL","MCO",0.00,0.00,
2019,3,2,"ATL","MCO",0.00,0.00,
2019,3,2,"ATL","MCO",,,
2019,3,2,"ATL","MCO",0.00,0.00,
2019,3,2,"ATL","MDW",0.00,0.00,
2019,3,2,"ATL","MDW",1.00,1.00,
[bitnami@debian:/tmp$ cat jan.csv febnh.csv marnh.csv > q1.csv
[bitnami@debian:/tmp$ head q1.csv
"YEAR","MONTH","DAY_OF_MONTH","ORIGIN","DEST","DEP_DEL15","ARR_DEL15",
2019,1,16,"BOS","PHL",0.00,0.00,
2019,1,17,"BOS","PHL",0.00,0.00,
2019,1,18,"BOS","PHL",0.00,0.00,
2019,1,19,"BOS","PHL",0.00,0.00,
2019,1,20,"BOS","PHL",,,
2019,1,21,"BOS","PHL",,,
2019,1,22,"BOS","PHL",,,
2019,1,23,"BOS","PHL",0.00,0.00,
2019,1,24,"BOS","PHL",0.00,0.00,
[bitnami@debian:/tmp$ tail q1.csv
2019,3,6,"ORD","CLT",0.00,0.00,
2019,3,6,"CVG","ORD",0.00,0.00,
2019,3,6,"DTW","ORD",0.00,0.00,
2019,3,6,"DTW","EWR",1.00,1.00,
2019,3,6,"OMA","EWR",1.00,1.00,
2019,3,6,"DTW","EWR",0.00,0.00,
2019,3,7,"ATL","ORD",0.00,0.00,
2019,3,7,"CLT","EWR",0.00,0.00,
2019,3,7,"ABQ","DEN",0.00,0.00,
2019,3,7,"CVG","ORD",0.00,0.00,
bitnami@debian:/tmp$
```

```

0: jdbc:hive2://localhost:10000/> create table q1_data (YEAR int, MONTH int, DAY int, AIRLINE_ID int, ORIGIN string, DEST string, DEL15 string,
ring, ARR15 string, cancelled string) row format delimited fields terminated by ',' stored as textfile;
INFO : Compiling command(queryId=hadoop_20211216050936_10c76385-ad48-4d3b-a2dc-ea5575f05b09): create table q1_data (YEAR int, MONTH int,
DAY int, AIRLINE_ID int, ORIGIN string, DEST string, DEL15 string, ARR15 string, cancelled string) row format delimited fields terminated
by ',' stored as textfile
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216050936_10c76385-ad48-4d3b-a2dc-ea5575f05b09); Time taken: 0.186 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216050936_10c76385-ad48-4d3b-a2dc-ea5575f05b09): create table q1_data (YEAR int, MONTH int,
DAY int, AIRLINE_ID int, ORIGIN string, DEST string, DEL15 string, ARR15 string, cancelled string) row format delimited fields terminated
by ',' stored as textfile
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216050936_10c76385-ad48-4d3b-a2dc-ea5575f05b09); Time taken: 1.839 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
No rows affected (2.058 seconds)
0: jdbc:hive2://localhost:10000/> describe q1_data;
[. . . . . > describe q1_data;
Error: Error while compiling statement: FAILED: ParseException line 2:0 cannot recognize input near 'q1_data' 'describe' 'q1_data' in tabl
e name (state=42000,code=40000)
0: jdbc:hive2://localhost:10000/> describe q1_data;
INFO : Compiling command(queryId=hadoop_20211216051121_b140ed61-f7b9-44e6-9217-0673132f06e3): describe q1_data
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:col_name, type:string, comment:from deserializer), FieldSchema(name:d
ata_type, type:string, comment:from deserializer), FieldSchema(name:comment, type:string, comment:from deserializer)], properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216051121_b140ed61-f7b9-44e6-9217-0673132f06e3); Time taken: 0.83 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216051121_b140ed61-f7b9-44e6-9217-0673132f06e3): describe q1_data
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216051121_b140ed61-f7b9-44e6-9217-0673132f06e3); Time taken: 0.134 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| year      | int       |         |
| month     | int       |         |
| day       | int       |         |
| airline_id | int       |         |
| origin    | string    |         |
| dest      | string    |         |
| del15     | string    |         |
| arr15     | string    |         |
| cancelled | string    |         |
+-----+-----+-----+
9 rows selected (1.03 seconds)

```

```

0: jdbc:hive2://localhost:10000/> create table q1_data (YEAR int, MONTH int, DAY int, AIRLINE_ID int, ORIGIN string, DEST string, DEL15 string,
ring, ARR15 string, cancelled string) row format delimited fields terminated by ',' stored as textfile;
INFO : Compiling command(queryId=hadoop_20211216050936_10c76385-ad48-4d3b-a2dc-ea5575f05b09): create table q1_data (YEAR int, MONTH int,
DAY int, AIRLINE_ID int, ORIGIN string, DEST string, DEL15 string, ARR15 string, cancelled string) row format delimited fields terminated
by ',' stored as textfile
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216050936_10c76385-ad48-4d3b-a2dc-ea5575f05b09); Time taken: 0.186 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216050936_10c76385-ad48-4d3b-a2dc-ea5575f05b09): create table q1_data (YEAR int, MONTH int,
DAY int, AIRLINE_ID int, ORIGIN string, DEST string, DEL15 string, ARR15 string, cancelled string) row format delimited fields terminated
by ',' stored as textfile
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216050936_10c76385-ad48-4d3b-a2dc-ea5575f05b09); Time taken: 1.839 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
No rows affected (2.058 seconds)
0: jdbc:hive2://localhost:10000/> describe q1_data;
[. . . . . > describe q1_data;
Error: Error while compiling statement: FAILED: ParseException line 2:0 cannot recognize input near 'q1_data' 'describe' 'q1_data' in tabl
e name (state=42000,code=40000)
0: jdbc:hive2://localhost:10000/> describe q1_data;
INFO : Compiling command(queryId=hadoop_20211216051121_b140ed61-f7b9-44e6-9217-0673132f06e3): describe q1_data
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:col_name, type:string, comment:from deserializer), FieldSchema(name:d
ata_type, type:string, comment:from deserializer), FieldSchema(name:comment, type:string, comment:from deserializer)], properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216051121_b140ed61-f7b9-44e6-9217-0673132f06e3); Time taken: 0.83 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216051121_b140ed61-f7b9-44e6-9217-0673132f06e3): describe q1_data
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216051121_b140ed61-f7b9-44e6-9217-0673132f06e3); Time taken: 0.134 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| year      | int       |         |
| month     | int       |         |
| day       | int       |         |
| airline_id | int       |         |
| origin    | string    |         |
| dest      | string    |         |
| del15     | string    |         |
| arr15     | string    |         |
| cancelled | string    |         |
+-----+-----+-----+
9 rows selected (1.03 seconds)

```

```

0: jdbc:hive2://localhost:10000/> create table q1_table (YEAR int, MONTH int, DAY int, ORIGIN string, DEST string, DEP_DEL15 string) row form
at delimited fields terminated by ',' stored as textfile;
INFO : Compiling command(queryId=hadoop_20211216053231_526dbe65-9784-4be1-9a62-f8bd5e77be11): create table q1_table (YEAR int, MONTH int, DA
Y int, ORIGIN string, DEST string, DEP_DEL15 string) row format delimited fields terminated by ',' stored as textfile
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216053231_526dbe65-9784-4be1-9a62-f8bd5e77be11); Time taken: 0.038 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216053231_526dbe65-9784-4be1-9a62-f8bd5e77be11): create table q1_table (YEAR int, MONTH int, DA
Y int, ORIGIN string, DEST string, DEP_DEL15 string) row format delimited fields terminated by ',' stored as textfile
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216053231_526dbe65-9784-4be1-9a62-f8bd5e77be11); Time taken: 0.169 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
No rows affected (0.226 seconds)
0: jdbc:hive2://localhost:10000/> load data local inpath '/tmp/q1nh.csv' overwrite into table q1_table;
INFO : Compiling command(queryId=hadoop_20211216053257_6587d656-e405-4d80-87d2-502e636b525a): load data local inpath '/tmp/q1nh.csv' overwri
te into table q1_table
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216053257_6587d656-e405-4d80-87d2-502e636b525a); Time taken: 0.095 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216053257_6587d656-e405-4d80-87d2-502e636b525a): load data local inpath '/tmp/q1nh.csv' overwri
te into table q1_table
INFO : Starting task [Stage-0:MOVE] in serial mode
INFO : Loading data to table default.q1_table from file:/tmp/q1nh.csv
INFO : Starting task [Stage-1:STATS] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216053257_6587d656-e405-4d80-87d2-502e636b525a); Time taken: 1.72 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
No rows affected (1.838 seconds)
0: jdbc:hive2://localhost:10000/> select * from q1_table limit 10;
INFO : Compiling command(queryId=hadoop_20211216053358_8ff3c899-142e-4173-91e9-b1e321e94888): select * from q1_table limit 10
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name=q1_table.year, type:int, comment:null), FieldSchema(name=q1_table.month,
type:int, comment:null), FieldSchema(name=q1_table.day, type:int, comment:null), FieldSchema(name=q1_table.origin, type:string, comment:null
), FieldSchema(name=q1_table.dest, type:string, comment:null), FieldSchema(name=q1_table.dep_del15, type:string, comment:null)], properties:n
ull)
INFO : Completed compiling command(queryId=hadoop_20211216053358_8ff3c899-142e-4173-91e9-b1e321e94888); Time taken: 0.365 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216053358_8ff3c899-142e-4173-91e9-b1e321e94888): select * from q1_table limit 10
INFO : Completed executing command(queryId=hadoop_20211216053358_8ff3c899-142e-4173-91e9-b1e321e94888); Time taken: 0.008 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager

```

q1_table.year	q1_table.month	q1_table.day	q1_table.origin	q1_table.dest	q1_table.dep_del15
2019	1	16	"BOS"	"PHL"	0.00
2019	1	17	"BOS"	"PHL"	0.00
2019	1	18	"BOS"	"PHL"	0.00
2019	1	19	"BOS"	"PHL"	0.00
2019	1	20	"BOS"	"PHL"	
2019	1	21	"BOS"	"PHL"	
2019	1	22	"BOS"	"PHL"	
2019	1	23	"BOS"	"PHL"	0.00
2019	1	24	"BOS"	"PHL"	0.00
2019	1	25	"BOS"	"PHL"	0.00

```

0: jdbc:hive2://localhost:10000/> create table q1_table_orc (YEAR int, MONTH int, DAY int, ORIGIN string, DEST string, DEP_DEL15 string) row
format delimited fields terminated by ',' stored as ORC;
INFO : Compiling command(queryId=hadoop_20211216053628_7dfe2711-ec14-45cb-8208-c72f52629617): create table q1_table_orc (YEAR int, MONTH int
, DAY int, ORIGIN string, DEST string, DEP_DEL15 string) row format delimited fields terminated by ',' stored as ORC
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216053628_7dfe2711-ec14-45cb-8208-c72f52629617); Time taken: 0.045 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216053628_7dfe2711-ec14-45cb-8208-c72f52629617): create table q1_table_orc (YEAR int, MONTH int
, DAY int, ORIGIN string, DEST string, DEP_DEL15 string) row format delimited fields terminated by ',' stored as ORC
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216053628_7dfe2711-ec14-45cb-8208-c72f52629617); Time taken: 0.131 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
No rows affected (0.198 seconds)
0: jdbc:hive2://localhost:10000/> create table q1_table_parquet (YEAR int, MONTH int, DAY int, ORIGIN string, DEST string, DEP_DEL15 string)
row format delimited fields terminated by ',' stored as parquet;
INFO : Compiling command(queryId=hadoop_20211216053714_a96a234c-e745-47b2-9454-32b11aa319bf): create table q1_table_parquet (YEAR int, MONTH
int, DAY int, ORIGIN string, DEST string, DEP_DEL15 string) row format delimited fields terminated by ',' stored as parquet
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:null, properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216053714_a96a234c-e745-47b2-9454-32b11aa319bf); Time taken: 0.049 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216053714_a96a234c-e745-47b2-9454-32b11aa319bf): create table q1_table_parquet (YEAR int, MONTH
int, DAY int, ORIGIN string, DEST string, DEP_DEL15 string) row format delimited fields terminated by ',' stored as parquet
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216053714_a96a234c-e745-47b2-9454-32b11aa319bf); Time taken: 0.11 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
No rows affected (0.192 seconds)
0: jdbc:hive2://localhost:10000/> show tables;
INFO : Compiling command(queryId=hadoop_20211216054221_807b9556-de9a-48a7-a421-454c75d46d94): show tables
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Semantic Analysis Completed (retrial = false)
INFO : Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:tab_name, type:string, comment:from deserializer)], properties:null)
INFO : Completed compiling command(queryId=hadoop_20211216054221_807b9556-de9a-48a7-a421-454c75d46d94); Time taken: 0.05 seconds
INFO : Concurrency mode is disabled, not creating a lock manager
INFO : Executing command(queryId=hadoop_20211216054221_807b9556-de9a-48a7-a421-454c75d46d94): show tables
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hadoop_20211216054221_807b9556-de9a-48a7-a421-454c75d46d94); Time taken: 0.029 seconds
INFO : OK
INFO : Concurrency mode is disabled, not creating a lock manager
+-----+
| tab_name |
+-----+
| q1_table |
| q1_table_orc |
| q1_table_parquet |
+-----+

```



```

> insert into table q1_table_orc select * from q1_table;
Automatically selecting local only mode for query
Query ID = root_20211216055517_7cf849f8-0266-4a0d-a344-d1f6a45d4be9
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2021-12-16 05:55:25,911 Stage-1 map = 0%, reduce = 0%
2021-12-16 05:55:41,318 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local1623538176_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://localhost:8020/user/hive/warehouse/q1_table_orc/.h
ive-staging_hive_2021-12-16_05-55-17_943_5102688540634637548-1/-ext-10000
Loading data to table default.q1_table_orc
MapReduce Jobs Launched:
Stage-Stage-1: HDFS Read: 113682032 HDFS Write: 3343457 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 25.899 seconds
hive> select * from q1_table_orc limit 10;
OK
2019 1 16 "BOS" "PHL" 0.00
2019 1 17 "BOS" "PHL" 0.00
2019 1 18 "BOS" "PHL" 0.00
2019 1 19 "BOS" "PHL" 0.00
2019 1 20 "BOS" "PHL"
2019 1 21 "BOS" "PHL"
2019 1 22 "BOS" "PHL"
2019 1 23 "BOS" "PHL" 0.00
2019 1 24 "BOS" "PHL" 0.00
2019 1 25 "BOS" "PHL" 0.00
Time taken: 0.401 seconds, Fetched: 10 row(s)
hive> insert into table q1_table_parquet select * from q1_table;
Automatically selecting local only mode for query
Query ID = root_20211216055710_25418f2d-5d5f-4705-9596-6432d7e203eb
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Job running in-process (local Hadoop)
2021-12-16 05:57:13,951 Stage-1 map = 0%, reduce = 0%
2021-12-16 05:57:30,172 Stage-1 map = 100%, reduce = 100%
Ended Job = job_local150877962_0002
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://localhost:8020/user/hive/warehouse/q1_table_parque
t/.hive-staging_hive_2021-12-16_05-57-10_503_6664825315255136668-1/-ext-10000
Loading data to table default.q1_table_parquet
MapReduce Jobs Launched:
Stage-Stage-1: HDFS Read: 230758294 HDFS Write: 9948353 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
Time taken: 20.849 seconds
hive> select * from q1_table_parquet limit 10;
OK
2019 1 16 "BOS" "PHL" 0.00
2019 1 17 "BOS" "PHL" 0.00
2019 1 18 "BOS" "PHL" 0.00
2019 1 19 "BOS" "PHL" 0.00
2019 1 20 "BOS" "PHL"
2019 1 21 "BOS" "PHL"
2019 1 22 "BOS" "PHL"
2019 1 23 "BOS" "PHL" 0.00
2019 1 24 "BOS" "PHL" 0.00
2019 1 25 "BOS" "PHL" 0.00
Time taken: 0.412 seconds, Fetched: 10 row(s)
hive> select * from q1_table limit 10;
OK
2019 1 16 "BOS" "PHL" 0.00
2019 1 17 "BOS" "PHL" 0.00
2019 1 18 "BOS" "PHL" 0.00
2019 1 19 "BOS" "PHL" 0.00
2019 1 20 "BOS" "PHL"
2019 1 21 "BOS" "PHL"
2019 1 22 "BOS" "PHL"
2019 1 23 "BOS" "PHL" 0.00
2019 1 24 "BOS" "PHL" 0.00
2019 1 25 "BOS" "PHL" 0.00
Time taken: 0.304 seconds, Fetched: 10 row(s)

```